

## The Brief

### Compensation:

We value your time and will compensate you with 150 Euros for completion of the task, regardless of the outcome.

### Context:

Our platform allows mobile apps to embed vertical video experiences (think TikTok-style feeds). Mobile SDKs fetch content and configuration from our backend APIs. Content teams use our CMS to manage videos, playlists, and display rules.

We're planning a new feature: **Personalised Video Feeds**. Instead of showing the same content to every user, we want to serve different videos based on user signals (watch history, engagement, demographic hints passed from the host app).

### Your Task:

Design and partially implement the backend system that would power personalised feeds.

### Non-negotiable constraints

- Scale: peak **3k RPS** to the feed endpoint (avg ~600 RPS).
- Latency: **p95 < 250 ms, p99 < 600 ms** for a **20-item** feed.
- Freshness: new content visible in **≤60 s**; user-signal updates can lag **≤5 min**.
- Privacy: SDK sends **hashed user ID**; no raw PII leaves our VNet; user-event retention **90 days**.
- Multi-tenant: **120 tenants**, each may override global ranking weights.
- Rollout: behind a **feature flag** with a safe **kill switch** to nonpersonalized feed.

### Deliverables

#### 1. System Design Document (primary deliverable)

- **Architecture diagram:** how personalisation fits API/CMS/SDK, where ranking happens, caching strategy, data flows for content and events.

- **Data model:** minimal tables/collections for user signals and content metadata; show keys and retention.
- **API contract:** endpoint(s) for the SDK (e.g., GET /v1/feed?limit=20), request/response schema, caching headers (ETag/Max-Age), and error semantics.
- **CMS configuration:** how content managers set rules (e.g., editorial boosts, maturity filters, per-tenant weights).
- **Trade-offs & decisions:** what you optimized for, what you deferred, and why - call out **reversible** steps.
- **Rollout & observability:** flags, fallback to non-personalized, the metrics/dashboards you'd ship first (adoption, p95/99, cold-start click-through).
- **What you'd do with more time:** e.g., bandits/ML, richer real-time signals, offline evals.

## 2. Working Prototype (secondary deliverable)

- A basic API endpoint that demonstrates your core personalisation logic
- Doesn't need to be production-ready - we're looking at your approach, code quality, and how you structure things
- Use whatever language/framework you're most productive in
- Include a README with setup instructions

## 3. Brief Write-up: How You Used AI

- What AI tools did you use during this task (if any)?
- What worked well? What didn't?
- How would you think about AI assistance for your team if you joined?

## What We're Evaluating

- **System thinking:** Can you design something that balances flexibility, simplicity, and real constraints?
- **Pragmatism:** Do you make sensible trade-offs for a small team with limited resources?
- **Code quality:** Is your prototype code clear, testable, and something you'd be comfortable having others build on?
- **Communication:** Can you explain technical decisions clearly in writing?
- **AI fluency:** Do you have a thoughtful, practical perspective on AI in development workflows?

## What We're NOT Looking For

- A fully working, deployed system
- Complex ML recommendation algorithms
- Perfect code coverage or enterprise patterns
- Any specific technology choice - use what makes you productive

